

A High Performance Implementation of the Escape Time Algorithm to generate the Mandelbrot Set

Parallel and high-performance computing

Emanuele Bugliarello

emanuele.bugliarello@epfl.ch

(SC-MA4)

June 29, 2017



Table of contents

- Scientific Background
- Implementations
- Going Parallel
- Results
- Conclusion

Scientific Background

Definition

The Mandelbrot set is the set of complex numbers c for which the function $f_c(z) = z^2 + c$ does not diverge when iterated from $z = 0$.

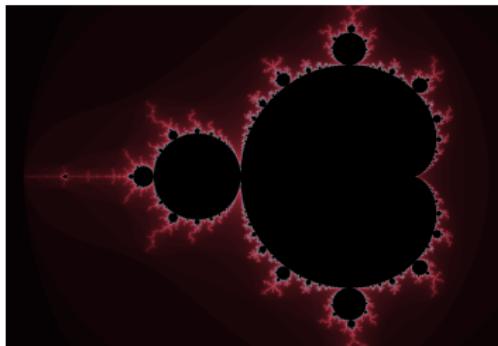
If we denote the Mandelbrot set by M , then by repeatedly applying the quadratic map

$$\begin{cases} z_0 = 0 \\ z_{n+1} = z_n + c, \end{cases}$$

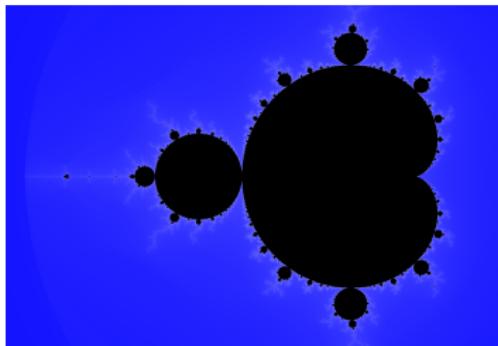
\forall complex number c , we have $c \in M \iff \limsup_{n \rightarrow \infty} |z_{n+1}| \leq 2$.

Generating Mandelbrot set images

- Sample complex numbers and iterate this function for each point c . If the function goes to infinity, the point belongs to the set.
- Each sampled point can be mapped on a 2D plane:
 - Treat its real and imaginary parts as image coordinates ($x + yi$)
 - Color this pixel according to how quickly $z_n^2 + c$ diverges



(a) Red-shaded



(b) Smooth, blue-shaded

Figure 1: Mandelbrot set drawings generated with our palettes.

The Escape Time Algorithm

Algorithm 1 Escape Time algorithm

```
1: for each pixel ( $P_x$ ,  $P_y$ ) on the screen do
2:    $x_0 \leftarrow$  scaled x coordinate of pixel
3:    $y_0 \leftarrow$  scaled y coordinate of pixel
4:    $x \leftarrow 0.0$ 
5:    $y \leftarrow 0.0$ 
6:    $iteration \leftarrow 0$ 
7:   while ( $x \times x + y \times y < 2 \times 2$  AND  $iteration < max\_iteration$ ) do
8:      $x_{temp} \leftarrow x \times x - y \times y + x_0$ 
9:      $y \leftarrow 2 \times x \times y + y_0$ 
10:     $x \leftarrow x_{temp}$ 
11:     $iteration \leftarrow iteration + 1$ 
12:     $color \leftarrow palette[iteration]$ 
13:     $plot(P_x, P_y, color)$ 
```

Where $z = x + iy$, $c = x_0 + iy_0$, $x = Re(z^2 + c)$ and $y = Im(z^2 + c)$.

Embarrassingly parallel problem: each pixel independent of any other.

Implementations

Implementations overview

Three main classes of implementations are developed:

- Serial versions.
- GPU versions using the CUDA parallel computing platform.
- Hybrid versions using both CUDA and MPI (including MPI-IO).

The application is coded in C.

Code fully debugged and profiled using `gdb` and `gprof`.

`Valgrind` used to make sure no memory leaks are possible.

Serial Implementations (1)

- naive.
- **opti1**: Exploit symmetry with respect to the x axis.
- **opti2**: Skip calculations for the points lying within the cardioid or in the period-2 bulb ($\approx 34\%$ of the total area).
- **opti3**: Finer-grained optimizations: reduce to 3 multiplications per loop (minimum) & use fast array indexing in plotting.

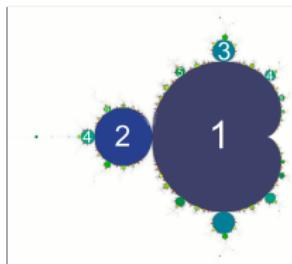


Figure 2: Periods of hyperbolic components.

Optimization level	Execution time [s]
-00	422.91
-01	230.99
-02	208.96
-03	208.49
-03 -ftree-vectorize	205.41

Table 1: Execution times of **opti3** ($368.64M$ pixels, $\text{max_iteration}=10,000$) as a function of the optimization level.

Serial Implementations (2)

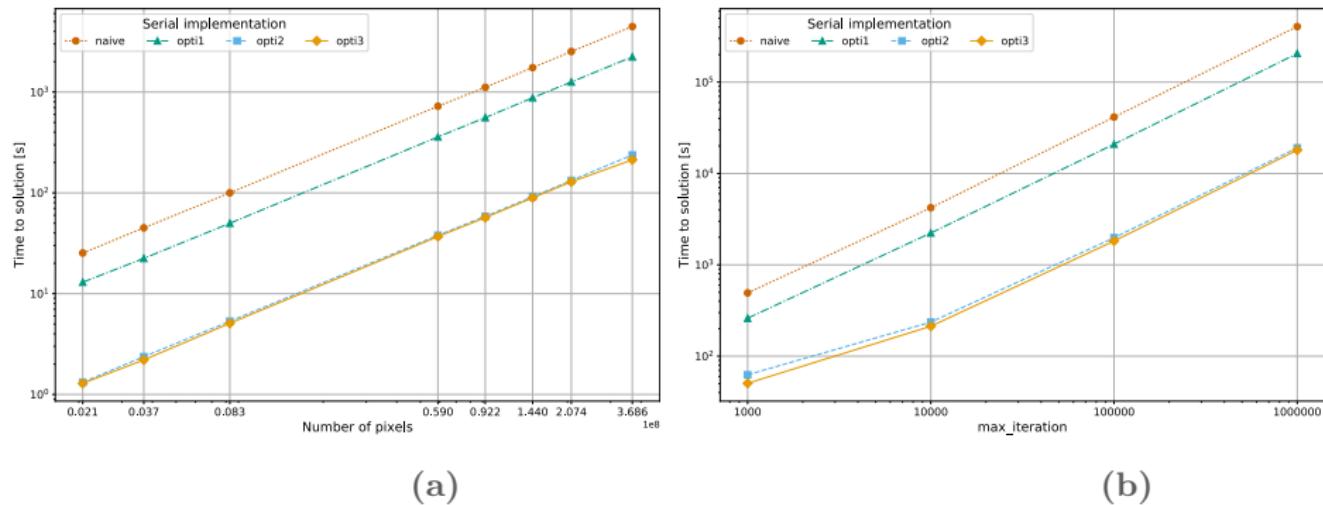


Figure 3: Time to solution for the different serial implementations with respect to (a) the number of pixels ($\text{max_iteration}=10,000$) and (b) max_iteration (368,640,000 pixels).

Going Parallel

CUDA Implementations (1)

- **opti3:** Parallelize escape time evaluation only.
 - Colors computed serially
 - 3 matrices sent on the PCIe bus
- **opti4:** All workload carried out by the GPU.
 - Only exchange RGB values of the bottom part on the PCIe bus
 - Use *Constant memory* to store constant values on the GPU
 - Code split into serial (.c) & parallel (.cu) parts; needs C wrapper
- **opti5:** Single-precision implementation of opti4.
 - Maximize instruction throughput
 - Cannot distinguish between double- and single-precision images

Application profiled with nvprof: compute-bound.

CUDA Implementations (2)

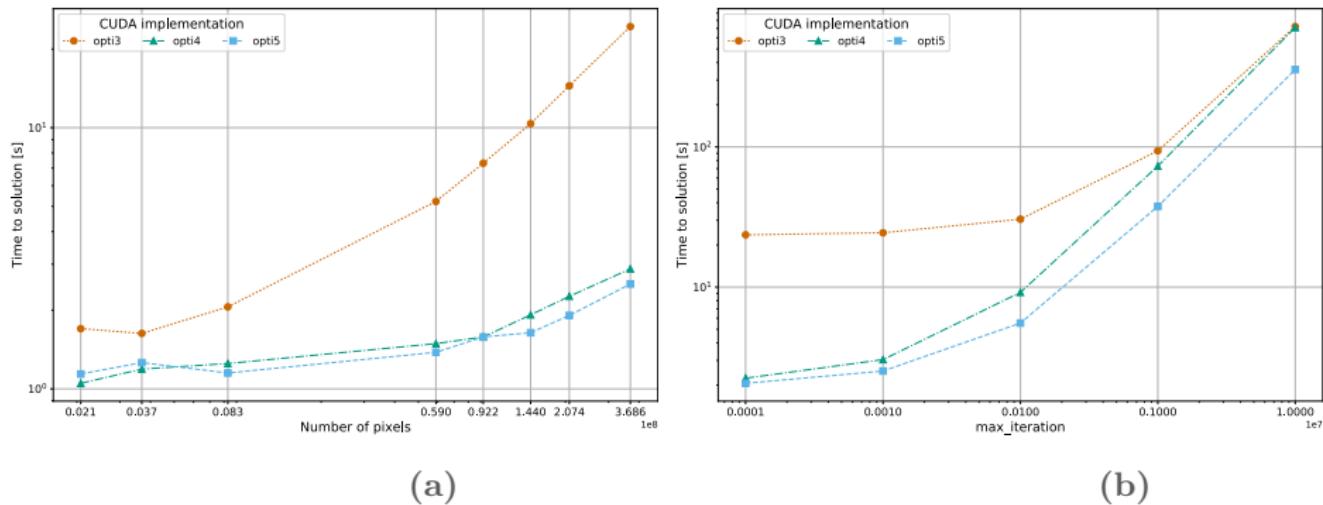


Figure 4: Time to solution for the CUDA implementations with respect to (a) the number of pixels ($\text{max_iteration}=10,000$) and (b) max_iteration (368,640,000 pixels) with block size 512×1 .

CUDA *block size* tuning for opti5

Block size:

- hyperparameter leading to performance improvements, increasing *occupancy* $\left(\frac{\text{active warps}}{\text{maximum active warps}} \right)$
- It depends on both the problem and the specifics of the GPU

Block size	Occupancy [%]		Block size	Execution time [s]	
	2560x1440	25600x14400		2560x1440	25600x14400
32x1	18.05	18.81	32x1	2.53	79.61
64x1	25.41	32.65	64x1	2.29	50.32
128x1	31.82	51.52	128x1	2.10	37.52
256x1	36.15	58.80	256x1	2.04	35.91
512x1	28.76	52.37	512x1	2.17	37.69
1024x1	8.39	25.50	1024x1	5.65	90.60
8x8	19.04	20.12	8x8	4.09	237.72
16x16	26.91	25.93	16x16	2.45	97.12
32x32	5.91	7.74	32x32	7.48	270.24

Table 2: Occupancy and execution time for two image sizes as the block size varies for max_iteration=10,000. Evaluated using nvprof.

Hybrid Implementations (1)

Distribute the computations over different CPUs and GPUs with MPI.

- Image split to the 2 processors in a node, each uses 1 GPU.
 - Non-blocking, point-to-point communications where rank 0-processor receives the image parts out-of-order ¹
 - Parallel I/O
- Image split to the 2 processors in a node, each uses 2 GPUs.
Highest performance: all the CUDA functions are asynchronous ²
 - Non-blocking, point-to-point communications where rank 0-processor receives the image parts out-of-order ¹
 - Parallel I/O

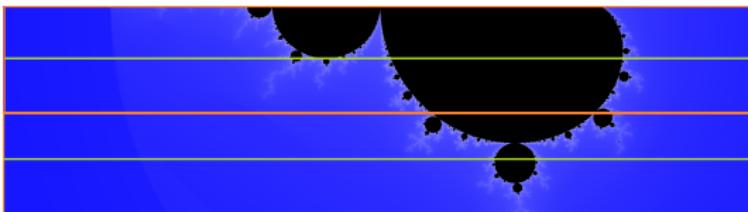
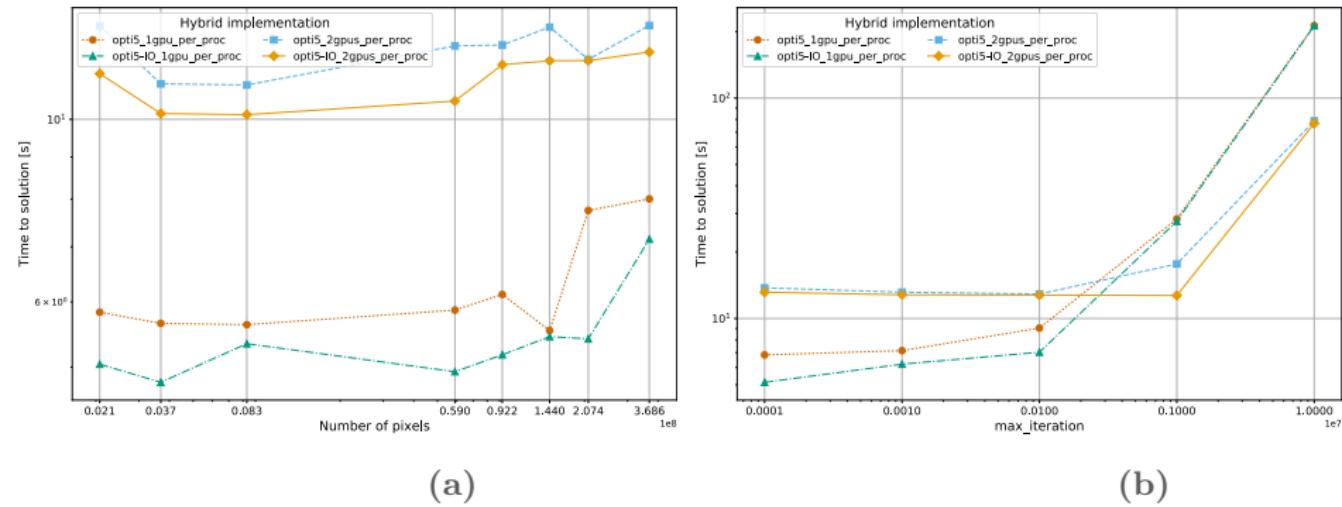


Figure 5: Image split to CPUs (orange) and GPUs (orange or green).

¹No `MPI_Barrier()` used.

²No `__syncthreads()` used.

Hybrid Implementations (2)



(a)

(b)

Figure 6: Time to solution for the hybrid implementations with respect to (a) the number of pixels ($\text{max_iteration}=10,000$) and (b) max_iteration (368,640,000 pixels) with block size 256×1 .

Results

Strong Scaling

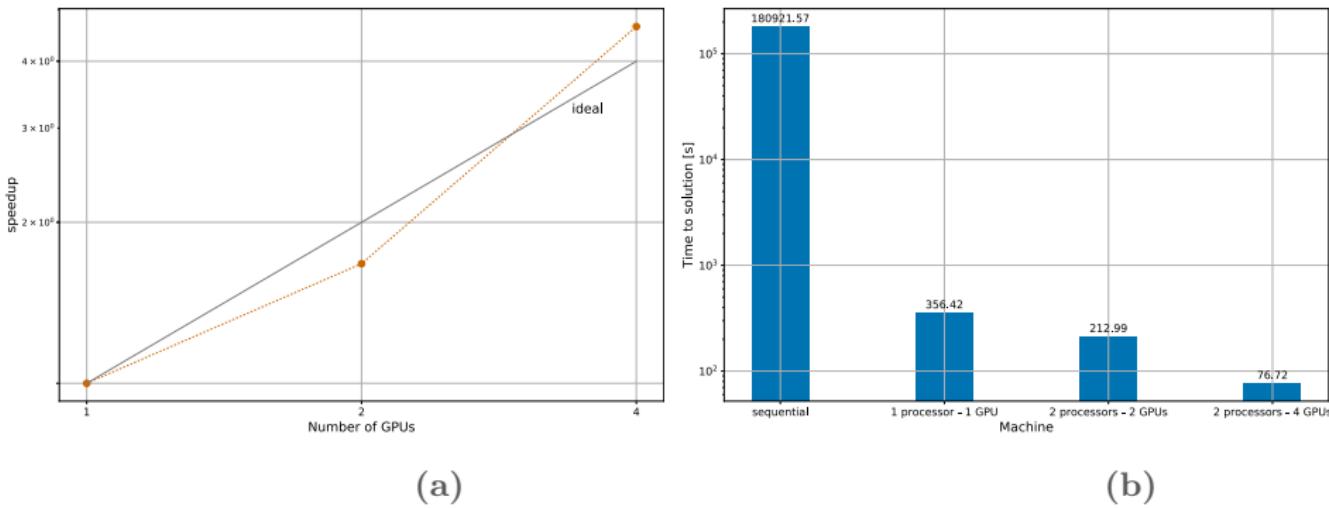


Figure 7: Strong scaling as (a) the speedup for an increasing number of GPUs (via MPI) and (b) the reduced time to solution of the same problem (368.64M-pixel image, $\text{max_iteration}=10 \times 10^6$) for the different versions.

Weak Scaling

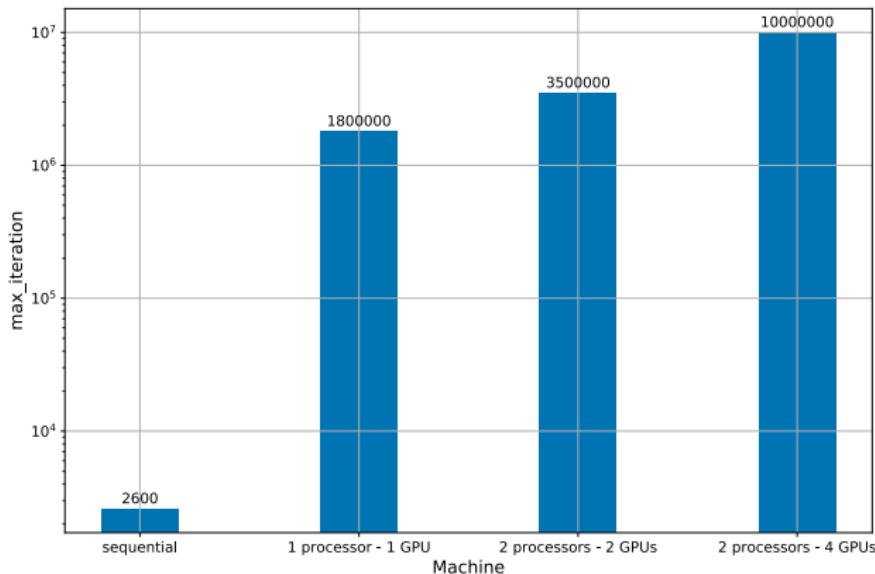


Figure 8: Weak scaling as a function of the problem size that can be solved in the same amount of time (76 s) by the different versions.

Conclusion

Conclusion & Future Work

Conclusion

- Implement and optimize `serial`, CUDA and CUDA+MPI applications
- Gain experience with CUDA and MPI in a supercomputer
- Acquire familiarity with profiling tools
- Investigate different technologies
 - Non-blocking communications and parallel I/O in MPI
 - Asynchronous functions in CUDA

Future Work

- Experiment with 2D grids
- Investigate the impact of external libraries
- Explore different algorithms

Conclusion & Future Work

Conclusion

- Implement and optimize `serial`, CUDA and CUDA+MPI applications
- Gain experience with CUDA and MPI in a supercomputer
- Acquire familiarity with profiling tools
- Investigate different technologies
 - Non-blocking communications and parallel I/O in MPI
 - Asynchronous functions in CUDA

Future Work

- Experiment with 2D grids
- Investigate the impact of external libraries
- Explore different algorithms

Thank you